

Sanjit Kumar

217-819-6384 | sanjitk3@illinois.edu | sanjitkumar.me | github.com/sanjitk7 | www.linkedin.com/in/sanjitk7 | San Jose, CA

EDUCATION:

University of Illinois Urbana-Champaign, *Master of Computer Science*

GPA: 3.91/4.00 | Aug 2022–May 2024

Relevant Coursework: Distributed Systems, OS Design, Fault Tolerant Data Center Systems, ML and Data Systems, Artificial Intelligence, Data Mining, Advanced Software Engineering Topics, Interactive Computer Graphics.

Vellore Institute of Technology, *Bachelor of Technology in Computer Science and Engineering*

GPA: 3.85/4.00 | Jul 2018–May 2022

Relevant Coursework: Data Structures and Algorithms, Database Management Systems, Object Oriented Programming, Web Programming, Computer Networks, Graph Theory and Social Networks, Parallel Programming.

SKILLS:

- **Programming Languages**: Python, C++, Java, C, Javascript, Go
- **Databases**: MySQL, MongoDB, Redis, Elasticsearch
- **Web Technologies**: ReactJS, Node.js, Express.js, PHP
- **Other Technologies**: Apache Kafka, Hadoop, Kibana, Docker, AWS, Linux Kernel, Bash, Data Plane Development Kit (DPDK), OpenGL, WebGL, PyTorch, Grafana, JavaFX.

EXPERIENCE:

University of Illinois Urbana Champaign, *Graduate Teaching Assistant*

C++, Go | Jan 2024 – May 2024 | Champaign, IL

- Worked with Dr. Radhika Mittal on the course [CS425: Distributed Systems](#) to provide industry relevant lessons for 200+ graduate/undergraduate students. Held bi-weekly office hours and assessed homework and projects.
- Mentored students to build distributed & networked application projects (reliable multicasting, distributed transactions, raft).

Aviz Networks Inc, *Fall Software Engineer Intern*

C, Python, Redis | Sep 2023 – Nov 2023 | San Jose, CA

- Developed and integrated a UI, REST API, and Redis DB Cache for a C-based Network Packet Collector for GTP Correlation.
- Worked with kernel-bypass strategies via Data Plane Development Kit to achieve packet processing speed for datacenter networks.
- Integrated an Apache Kafka module using Kafka-C drivers into the packet data processing pipeline, enhancing reliability.
- Optimized and performance-tuned the packet collector pipeline via payload batching. Achieved a 3x improvement in throughput and 50% reduction in latency.
- Automated the configuration, build and execution of the packet collector with Bash scripts and Python.

Aviz Networks Inc, *Summer Software Engineer Intern*

C++, Python | May 2023 – Aug 2023 | San Jose, CA

- Designed and developed a C++ based scalable network packet analyzer system. System tapped packets from high throughput network traffic from data centers to extract and stream metadata information.
- Leveraged scalable event processing systems like Kafka and Elasticsearch for data pipelining and downstream analytics.
- Benchmarked scalability on physical network devices with software-based (Scapy) and hardware-based (Ixia) load generators.
- Used a Kafka consumer to integrate a REST API with the system for selective packet capture and data sink integrations.

WebKnot Technologies, *Full Stack Developer Intern*

ReactJS, Node.js | Nov 2020 – Dec 2020 | Bangalore, India

- Integrated custom Tensorflow-based object detection models with Shinobi, an open-source CCTV framework.
- Developed UI with ReactJS and wrote REST API endpoints with node.js and Express.js for two different MERN stack applications.

PROJECTS:

Distributed ML Inference Job Scheduler

[GitHub](#) | Java, Python, PyTorch

- Designed a distributed job scheduler system for ML inference tasks built on top of 10 Linux VMs from scratch using Java and Python.
- Uses a real-time work scheduling algorithm to optimize query rate for ResNet and ImageNet classification tasks.
- Includes a distributed data logging service, distributed group membership protocol and failure detector, a distributed files system.

Page Fault Profiler

[GitHub](#) | C, Bash

- A page fault profiler in the linux kernel space to profile major/minor page faults and cpu run times of user processes via a virtual shared memory buffer. Implemented regular sampling via delayed work queue and user process registrations via proc interface.

Modified Banker's Algorithm with Software Interrupts

[GitHub](#) | C, Bash

- Enhanced the Banker's Algorithm to simulate standard operations but also manage OS software interrupts with a circular queue. Implemented the algorithm using a multi-threaded approach with the POSIX library in C.

WildSprint: Wildlife Fundraising Platform (*1st Place Ethereum Track – DevSpace '20 Hack*)

[GitHub](#) | Node.js, ReactJS, Solidity

- Developed a MERN stack fundraising platform to raise funds via cryptocurrency (wrote smart contracts to receive ETH coin).
- Built a live stream feature through [Dyte SDK](#) integration. Livestreams of wildlife in national parks to incentivize donations.

RESEARCH EXPERIENCE:

Internet of Things Security: Attacks, Solutions, Strengths and Limitations

[Link](#) | Sep 2021

International Conference on Artificial Intelligence and Machine Vision, IEEE

- Presented a comparative analysis of benchmarks between latest security frameworks in then-recent IoT literature while advised by Dr. Anil Kumar Kakelli.
- Categorically classified and critiqued existing IoT security frameworks based on their approaches to address the threat of malignant nodes in heterogeneous device networks and general strengths/limitations.